



RTX Pro 6000 (Blackwell) vs DGX Spark (GB10)

FP8 • sglang • 2048→2048
Consistent Colors Across Batches


Prefill Throughput (Tokens/sec) — Batch Size 1




Decode Throughput (Tokens/sec) — Batch Size 1




Estimated E2E Latency (s) — 2k in / 2k out — Batch Size 1




Prefill Throughput (Tokens/sec) — Batch Size 2




Decode Throughput (Tokens/sec) — Batch Size 2




Estimated E2E Latency (s) — 2k in / 2k out — Batch Size 2




Prefill Throughput (Tokens/sec) — Batch Size 4




Decode Throughput (Tokens/sec) — Batch Size 4




Estimated E2E Latency (s) — 2k in / 2k out — Batch Size 4




Prefill Throughput (Tokens/sec) — Batch Size 8




Decode Throughput (Tokens/sec) — Batch Size 8




Estimated E2E Latency (s) — 2k in / 2k out — Batch Size 8




Prefill Throughput (Tokens/sec) — Batch Size 16




Decode Throughput (Tokens/sec) — Batch Size 16




Estimated E2E Latency (s) — 2k in / 2k out — Batch Size 16



Prefill Throughput (Tokens/sec) — Batch Size 32



Decode Throughput (Tokens/sec) — Batch Size 32



Estimated E2E Latency (s) — 2k in / 2k out — Batch Size 32

